

Outline

[1. Establishing Goals](#)

[2. Introduction to Dataset](#)

[3. Exploratory Data Analysis](#)

[4. Limit Data to Scope](#)

[5. Supervised Feature Generation](#)

- [5.1 Common Bag of Words](#)
- [5.2 Turn Common Words into Features](#)
- [5.3 Clustering on BOW](#)
- [5.4 Classification on BOW](#)

[6. Unsupervised Feature Generation](#)

- [6.1 Latent Semantic Analysis](#)
- [6.2 Clustering on LSA \(BOW Content\)](#)
- [6.3 Classification on LSA \(BOW Content\)](#)
- [6.4 Clustering on LSA \(All Content\)](#)
- [6.5 Classification on LSA \(All Content\)](#)

[7. Choosing Model](#)

- [7.1 Comparing Scores](#)
- [7.2 Sorting by Test Accuracy](#)
- [7.3 Winner](#)

1. Establishing Goals

In this project I'll attempt to build models to correctly predict the author of a given article. The scope will be limited to 10 authors. The techniques I'll compare will include `Bag-of-Words` VS `Latent Semantic Analysis` for feature-generation, and `Clustering` VS `Supervised Learning` for classification. I'll also experiment with different sample sizes, as feature-generation can be very sensitive to high dimensionality.

2. Introduction to DataSet

From: <https://www.kaggle.com/snapcrack/all-the-news>

This dataset contains news articles scraped from various publications, labeled by publication and author name, as well as date and title.

The original source on `kaggle.com` contains three `.csv` files. Across the three, there are over 140,000 articles from a total of 15 publications.

The dataset used here is only the first of those three files, which contains about a third of all the data at roughly `280MB`. This is more than enough data for the goals of this project.

3. Exploratory Data Analysis

Let's get a quick overview of the data available.

```
1  # General-purpose Libraries
2  import numpy as np
3  import pandas as pd
4  import scipy
5  import sklearn
6  import spacy
7  import matplotlib.pyplot as plt
8  import seaborn as sns
9  import re
10 from collections import Counter
11 import spacy
12 from time import time
13 %matplotlib inline
14
15 # Tools for processing data
16 from sklearn.pipeline import make_pipeline
17 from sklearn.preprocessing import Normalizer
18 from sklearn.decomposition import TruncatedSVD, PCA
19 from sklearn.model_selection import cross_val_score, train_test_split, GridSearchCV
```

```

20 from sklearn.metrics import accuracy_score, recall_score, classification_report, confusion_matrix, make_scorer,
   adjusted_rand_score, silhouette_score, homogeneity_score, normalized_mutual_info_score
21 # Classifiers, supervised and unsupervised
22 from sklearn import ensemble
23 from sklearn.linear_model import LogisticRegression
24 from sklearn.svm import SVC
25 from sklearn.feature_extraction.text import TfidfVectorizer
26 from sklearn.cluster import KMeans
27 from sklearn.cluster import MeanShift, estimate_bandwidth
28 from sklearn.cluster import SpectralClustering
29 from sklearn.cluster import AffinityPropagation
30
31 import warnings
32 warnings.filterwarnings("ignore")

```

```

1 # Read data into a DataFrame
2 data = pd.read_csv("articles1.csv")

```

```

1 # Preview the data
2 data.head(3)

```

	Unnamed: 0	id	title	publication	author	date	year	month	url	content
0	0	17283	House Republicans Fret About Winning Their Hea...	New York Times	Carl Hulse	2016-12-31	2016.0	12.0	NaN	WASHINGTON — Congressional Republicans have...
1	1	17284	Rift Between Officers and Residents as Killing...	New York Times	Benjamin Mueller and Al Baker	2017-06-19	2017.0	6.0	NaN	After the bullet shells get counted, the blood...
2	2	17285	Tyrus Wong, 'Bambi' Artist Thwarted by Racial ...	New York Times	Margalit Fox	2017-01-06	2017.0	1.0	NaN	When Walt Disney's "Bambi" opened in 1942, cri...

Checking for Missing Data

- The content feature is complete. That's the most important thing. Some author names are missing. We'll make sure to choose 10 properly labeled.

```

1 data.info()

```

```

1 <class 'pandas.core.frame.DataFrame'>
2 RangeIndex: 50000 entries, 0 to 49999
3 Data columns (total 10 columns):
4 Unnamed: 0    50000 non-null int64
5 id           50000 non-null int64
6 title        50000 non-null object
7 publication   50000 non-null object
8 author       43694 non-null object
9 date         50000 non-null object
10 year        50000 non-null float64
11 month       50000 non-null float64
12 url         0 non-null float64
13 content     50000 non-null object
14 dtypes: float64(3), int64(2), object(5)
15 memory usage: 3.8+ MB

```

Length of Articles

- In terms of number of characters, the average article has less than 4,000 letters.

```



1 lengths = pd.Series([len(x) for x in data.content])
2 print('Statistical Summary of Article Lengths')
3 print(lengths.describe())
4
5 sns.distplot(lengths,kde=False)
6 plt.title('Distribution of Article Lengths (All)')
7 plt.show()
8 sns.distplot(lengths[lengths<10000],kde=False)
9 plt.title('Distribution of Articles Lengths < 10,000 Characters')
10 plt.show()

```

```

1 Statistical Summary of Article Lengths
2 count      50000.00000
3 mean       3853.4537
4 std        3875.9117
5 min         1.0000
6 25%        1682.0000
7 50%        2853.0000
8 75%        5045.0000
9 max       149346.0000
10 dtype: float64

```



4. Limit Data to Scope

[Back to Outline](#)

Here I'll pick the 10 authors whose names I'll predict based on their content. This selection will remain the same for all the methods I'll compare.

Since we only need 10 authors, I'll get the first 10 authors whose article-count is greater than X. 100 articles per author is a good number because more would take terribly long when fit to classifiers after `TF-IDF`. At the same time, `Bag-of-Words` is the slowest. However, for that I'll limit to 50 of these articles per author.

```

1 # First ten authors with more than X articles
2 print(data.author.value_counts()[data.author.value_counts().>100][-10:])

```

```

1 Scott Davis      119
2 Eugene Scott    118
3 Laura Smith-Spark 115
4 Julie Bort      110
5 Raheem Kassam   110
6 Jeremy Berke    109
7 Eli Watkins     106
8 Oliver Darcy    104
9 Daniella Diaz   104
10 Cartel Chronicles 102
11 Name: author, dtype: int64

```

```

1 # Make a DataFrame with articles by our chosen authors
2 # Include author names and article titles.
3
4 # Make a list of the 10 chosen author names
5 names = data.author.value_counts()[data.author.value_counts(>100)[-10:].index.tolist()
6
7 # DataFrame for articles of all chosen authors
8 authors_data = pd.DataFrame()
9 for name in names:
10     # Select each author's data
11     articles = data[data.author==name][:100][['title', 'content', 'author']]
12     # Append it to the DataFrame
13     authors_data = authors_data.append(articles)
14
15 authors_data = authors_data.reset_index().drop('index',1)
16
17 authors_data.head()

```

	title	content	author
0	A scramble for quarterbacks in the 2016 NFL Dr...	''' Two NFL teams enter the postseason st...	Scott Davis
1	Rio's Olympic Stadium has reportedly turned in...	''' As is the case with many Rio's Marac...	Scott Davis
2	The Grizzlies gambled on Chandler Parsons with...	''' Even in an NBA era with a rising sala...	Scott Davis
3	Kevin Love had some simple advice for the Cavs...	''' The Cleveland Cavaliers are with the...	Scott Davis
4	Aaron Rodgers completes another ridiculous Hai...	' After a slow start to their Wild Card game...	Scott Davis

```

1 # Look for duplicates
2 print('Number of articles:',authors_data.shape[0])
3 print('Unique articles:',len(np.unique(authors_data.index)))
4
5 # Number of authors
6 print('Unique authors:',len(np.unique(authors_data.author)))
7 print('')
8 print('Articles by author:\n')
9
10 # Articles counts by author
11 print(authors_data.author.value_counts())

```

```



1 Number of articles: 1000
2 Unique articles: 1000
3 Unique authors: 10
4
5 Articles by author:
6
7 Oliver Darcy      100
8 Cartel Chronicles 100
9 Jeremy Berke     100
10 Daniella Diaz    100
11 Laura Smith-Spark 100
12 Eli Watkins      100
13 Eugene Scott     100
14 Raheem Kassam    100
15 Scott Davis      100
16 Julie Bort       100
17 Name: author, dtype: int64

```

Look at the Size of Articles Chosen

```
1 lengths = pd.Series([len(x) for x in authors_data.content])
2 print('Statistical Summary of Article Lengths')
3 print(lengths.describe())
4
5 sns.distplot(lengths,kde=False)
6 plt.title('Distribution of Article Lengths (All)')
7 plt.show()
8 sns.distplot(lengths[lengths<10000],kde=False)
9 plt.title('Distribution of Articles Lengths < 10,000 Characters')
10 plt.show()
```

```
1 Statistical Summary of Article Lengths
2 count      1000.000000
3 mean       3004.200000
4 std        2608.965556
5 min         106.000000
6 25%        1645.000000
7 50%        2356.500000
8 75%        3522.500000
9 max        33798.000000
10 dtype: float64
```



5. Supervised Feature Generation

[Back to Outline](#)

Bag of words is a list of the most common words of a given source of text. To identify each author, I'll create a bag of words containing the most-common words of all authors combined. This set later becomes the basis for feature engineering.

5.1 Common Bag of Words

- Here I'll extract the most-common 1000 words from each author's corpus, store them in a list, and then eliminate duplicates.

```
1 t0 = time()
2
3 # Load spacy NLP object
4 nlp = spacy.load('en')
5
6 # A list to store common words by all authors
7 common_words = []
8
```

```

9  # A dictionary to store the spacy_doc object of each author
10 authors_docs = {}
11
12 for name in names:
13     # Corpus is all the text written by that author
14     corpus = ""
15     # Grab all rows of current author, along the 'content' column
16     author_content = authors_data.loc[authors_data.author==name, 'content']
17
18     # Merge all articles in to the author's corpus
19     for article in author_content:
20         corpus = corpus + article
21     # Let Spacy parse the author's body of text
22     doc = nlp(corpus)
23
24     # Store the doc in the dictionary
25     authors_docs[name] = doc
26
27     # Filter out punctuation and stop words.
28     lemmas = [token.lemma_ for token in doc
29                if not token.is_punct and not token.is_stop]
30
31     # Return the most common words of that author's corpus.
32     bow = [item[0] for item in Counter(lemmas).most_common(1000)]
33
34     # Add them to the list of words by all authors.
35     for word in bow:
36         common_words.append(word)
37
38     # Eliminate duplicates
39     common_words = set(common_words)
40
41 print('Total number of common words:', len(common_words))
42 print("done in %.3fs" % (time() - t0))

```

```

1  Total number of common words: 3658
2  done in 71.345s

```

- From a theoretical total of 10,000 common-words, (1,000 from 10 authors) 3,405 were unique. So roughly a third of all words used by each author is actually part of their unique style.

```

1  # Let's see our 10 authors in the dictionary
2  lengths = []
3  for k,v in authors_docs.items():
4      print(k, 'corpus contains', len(v), ' words.')
5      lengths.append(len(v))

```

```

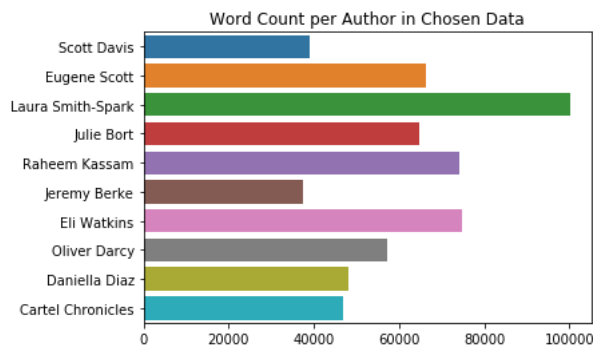
1  Scott Davis corpus contains 39034 words.
2  Eugene Scott corpus contains 66174 words.
3  Laura Smith-Spark corpus contains 100065 words.
4  Julie Bort corpus contains 64793 words.
5  Raheem Kassam corpus contains 74184 words.
6  Jeremy Berke corpus contains 37497 words.
7  Eli Watkins corpus contains 74679 words.
8  Oliver Darcy corpus contains 57252 words.
9  Daniella Diaz corpus contains 48219 words.
10 Cartel Chronicles corpus contains 46935 words.

```

```

1  sns.barplot(x=lengths, y=names, orient='h')
2  plt.title('Word Count per Author in Chosen Data')
3  plt.show()

```



5.2 Turn Common Words into Features

Approach

Due to the curse of dimensionality, doing this step with all our 1,000 articles would take prohibitively long. (10 authors * 100 articles/ea = 1,000 articles) At 30 seconds per article, my personal machine would need 8.5 hours of processing. Therefore I'll limit this part to 50 articles per author. This should still convey enough information for a decent predictive model.

About 'Common Bag of Words'

This technique consists of creating a feature out of each common word and then counting the number of times each common word appears in each article. Each cell will represent the number of times the lemma of the given column appears in the article of the current row. We have over 3,000 common words, and will be using 500 articles total. (50 per author) Plus each article may have a varying number of words in it. That's a lot of text to compare and count.

```
1 # check for lower case words
2 common_words = pd.Series(pd.DataFrame(columns=common_words).columns)
3 print('Count of all common_words:', len(common_words))
4 print('Count of lowercase common_words:', np.sum([word.islower() for word in common_words]))
5
6 # Turn all common_words into lower case
7 common_words = [word.lower() for word in common_words]
8 print('Count of lowercase common_words (After Conversion):', np.sum([word.islower() for word in common_words]))
```

```
1 Count of all common_words: 3658
2 Count of lowercase common_words: 2352
3 Count of lowercase common_words (After Conversion): 3579
```

- Notice that after converting to lowercase the total number of lowercase words still isn't the same as the total. This means there are around 100 non alphabetic words inside our bag. This is probably made up of numbers and words with punctuations within.

```
1 #We must remove these in to avoid conflicts with existing columns.
2
3 if 'author' in common_words:
4     common_words.remove('author')
5 if 'title' in common_words:
6     common_words.remove('title')
7 if 'content' in common_words:
8     common_words.remove('content')
9
```

```
1 # Count the number of times a common_word appears in each article
2 # (about 3Hrs processing)
3
4 bow_counts = pd.DataFrame()
5 for name in names:
6     # Select X articles of that author
7     articles = authors_data.loc[authors_data.author==name,:][:50]
8     bow_counts = bow_counts.append(articles)
9 bow_counts = bow_counts.reset_index().drop('index',1)
10
11 # Use common_words as the columns of a temporary DataFrame
12 df = pd.DataFrame(columns=common_words)
13
14 # Join BOW features with the author's content
15 bow_counts = bow_counts.join(df)
16
17 # Initialize rows with zeroes
18 bow_counts.loc[:, common_words] = 0
```

```

19
20 # Fill the DataFrame with counts of each feature in each article
21 t0 = time()
22 for i, article in enumerate(bow_counts.content):
23     doc = nlp(article)
24     for token in doc:
25         if token.lemma_.lower() in common_words:
26             bow_counts.loc[i, token.lemma_.lower()] += 1
27     # Print a message every X articles
28     if i % 50 == 0:
29         if time()-t0 < 3600: # if less than an hour in seconds
30             print("Article ", i, " done after ", (time()-t0)/60, ' minutes.')
31         else:
32             print("Article ", i, " done after ", (time()-t0)/60/60, ' hours.')

```

```

1 Article  0  done after  0.42328090270360313  minutes.
2 Article 50  done after 10.271135667959848  minutes.
3 Article 100 done after 24.337886516253153  minutes.
4 Article 150 done after 46.48636318047841  minutes.
5 Article 200 done after 1.094332391752137  hours.
6 Article 250 done after 1.3873678059710397  hours.
7 Article 300 done after 1.5501557772027121  hours.
8 Article 350 done after 1.805334949957  hours.
9 Article 400 done after 2.171408551865154  hours.
10 Article 450 done after 2.347395773132642  hours.

```

- This is the data that we can use to train clusters and classifiers. Each entry is an article, each column is a common word, and each cell is a count of the current common word in the current article.

```
1 bow_counts.head(3)
```

	title	content	author	becker	firm	dominate	russians	have	undermine	iran	...	hear	activity
0	A scramble for quarterbacks in the 2016 NFL Dr...	''' Two NFL teams enter the postseason st...	Scott Davis	0	0	0	0	12	0	0	...	0	0
1	Rio's Olympic Stadium has reportedly turned in...	''' As is the case with many Rio's Marac...	Scott Davis	0	0	0	0	6	0	0	...	0	0
2	The Grizzlies gambled on Chandler Parsons with...	''' Even in an NBA era with a rising sala...	Scott Davis	0	0	0	0	3	0	0	...	0	0

3 rows × 3658 columns

Optional:

- Store contents of `bow_counts`

```

1 # This saves the long-awaited data into a pickle file for easy recovery
2 #bow_counts.to_pickle('bow_counts')
3
4 # Read it back in with the following
5 #bow_counts = pd.read_pickle('bow_counts')

```

```

1 # Make sure we have 50 articles per author
2 bow_counts.author.value_counts()

```



```

1 | Eli Watkins      50
2 | Eugene Scott    50
3 | Oliver Darcy    50
4 | Raheem Kassam   50
5 | Cartel Chronicles 50
6 | Scott Davis     50
7 | Jeremy Berke    50
8 | Julie Bort      50
9 | Daniella Diaz   50
10 | Laura Smith-Spark 50
11 | Name: author, dtype: int64

```

5.3. Clustering on BOW

[Back to Outline](#)

- Before classifying, I'll start with clustering. Here I'll create clusters out of the BOW data and see if those clusters resemble the actual author's content. Clusters have no labels, but similar content tends to fall into the same clusters. Therefore in an ideal clustering solution, each author's articles would all fall into a single cluster.

```

1 | # Establish outcome and predictors
2 | y = bow_counts['author']
3 | X = bow_counts.drop(['content', 'author', 'title'], 1)
4 |
5 | X_train, X_test, y_train, y_test = train_test_split(X, y,
6 |                                                    test_size=0.24,
7 |                                                    random_state=0,
8 |                                                    stratify=y)

```

```

1 | # Make sure classes are balanced after train-test-split
2 | y_test.value_counts()

```

```

1 | Raheem Kassam      12
2 | Julie Bort         12
3 | Eli Watkins        12
4 | Eugene Scott       12
5 | Cartel Chronicles   12
6 | Scott Davis        12
7 | Jeremy Berke       12
8 | Laura Smith-Spark   12
9 | Oliver Darcy       12
10 | Daniella Diaz      12
11 | Name: author, dtype: int64

```

DataFrame to Store our Results

This `DataFrame` will hold results from all algorithms implemented ahead. For clustering algorithms, the `train/test` and `cross_val` columns will be left blank because clustering requires no train/test split. On the other hand, classifiers will indeed store their own `ARI`, `Homogeneity`, `Silhouette`, and `Mutual_Info` scores. `Features` will represent the method for feature-engineering, whether BOW or LSA. And the `n_train` column will represent the number of samples in the train size.

```

1 | # Store our results in a DataFrame
2 | metrics = ['Algorithm', 'n_train', 'Features', 'ARI', 'Homogeneity',
3 |           'Silhouette', 'Mutual_Info', 'Cross_Val', 'Train_Accuracy',
4 |           'Test_Accuracy']
5 | performance = pd.DataFrame(columns=metrics)

```

Approach to Clustering

In cluster analysis, there usually is no training or test data split. Because you do cluster analysis when you do not have labels, so you cannot “train”. Training is a concept from machine learning, and train-test splitting is used to avoid overfitting. But if you are not learning labels, you cannot overfit. Properly used cluster analysis is a knowledge discovery method. You want to discover some new structure in your data, not rediscover something that is already labeled.

5.3.1. Unsupervised Parameter Search Function

[Back to Outline](#)

- This function will find the parameters that produce the highest `Normalized Mutual Infomation` score from our clusters. This score is a good baseline from which to compare clustering VS classification because it correlates with good clutering as well as

higher accuracy scores.

- It'll print the relevant statistics as well as a contingency matrix of the result and lastly store our results in an external DataFrame.


```
1 # Function to quickly evaluate clustering solutions
2 def evaluate_clust(clust,params,features,i):
3     t0 = time()
4     print('\n','-'*40,'\n',clust.__class__.__name__,'\n','-'*40)
5
6     # Find best parameters based on scoring of choice
7     score = make_scorer(normalized_mutual_info_score)
8     search = GridSearchCV(clust,params,scoring=score,cv=3).fit(X,y)
9     print("Best parameters:",search.best_params_)
10    y_pred = search.best_estimator_.fit_predict(X)
11
12    ari = adjusted_rand_score(y, y_pred)
13    performance.loc[i,'ARI'] = ari
14    print("Adjusted Rand-Index: %.3f" % ari)
15
16    hom = homogeneity_score(y,y_pred)
17    performance.loc[i,'Homogeneity'] = hom
18    print("Homogeneity Score: %.3f" % hom)
19
20    sil = silhouette_score(X,y_pred)
21    performance.loc[i,'Silhouette'] = sil
22    print("Silhouette Score: %.3f" % sil)
23
24    nmi = normalized_mutual_info_score(y,y_pred)
25    performance.loc[i,'Mutual_Info'] = nmi
26    print("Normed Mutual-Info Score: %.3f" % nmi)
27
28    performance.loc[i,'n_train'] = len(X)
29    performance.loc[i,'Features'] = features
30    performance.loc[i,'Algorithm'] = clust.__class__.__name__
31
32    # Print contingency matrix
33    crosstab = pd.crosstab(y, y_pred)
34    plt.figure(figsize=(8,5))
35    sns.heatmap(crosstab, annot=True,fmt='d', cmap=plt.cm.copper)
36    plt.show()
37    print(time()-t0,"seconds.")
```

5.3.2. KMeans CBOW

[Back to Outline](#)

```
1 clust=KMeans()
2 params={
3     'n_clusters': np.arange(10,30,5),
4     'init': ['k-means++','random'],
5     'n_init':[10,20],
6     'precompute_distances':[True,False]
7 }
8 evaluate_clust(clust,params,features='BOW',i=0)
```

```
1 -----
2 KMeans
3 -----
4 Best parameters: {'init': 'random', 'n_clusters': 20, 'n_init': 20, 'precompute_distances': True}
5 Adjusted Rand-Index: 0.210
6 Homogeneity Score: 0.447
7 Silhouette Score: 0.050
8 Normed Mutual-Info Score: 0.428
```




```
1 169.47527551651 seconds.
```

5.3.3. Mean Shift CBOW

[Back to Outline](#)

```
1 #Declare and fit the model
2 clust = MeanShift()
3
4 params={}
5 evaluate_clust(clust,params,features='BOW',i=1)
```

```
1 -----
2 MeanShift
3 -----
4 Best parameters: {}
5 Adjusted Rand-Index: 0.002
6 Homogeneity Score: 0.101
7 Silhouette Score: 0.277
8 Normed Mutual-Info Score: 0.171
```



```
1 172.52272963523865 seconds.
```

- The above is a really bad solution. 30 clusters were created but most of our articles were assigned to the first cluster.

5.3.4. Affinity Propagation CBOW


[Back to Outline](#)

```
1 #Declare and fit the model.
2 clust = AffinityPropagation()
3
4 params = {
5     'damping': [.5, .7, .9],
6     'max_iter': [200, 500]
7 }
8 evaluate_clust(clust,params,features='BOW',i=2)
```

```

1 -----
2 AffinityPropagation
3 -----
4 Best parameters: {'damping': 0.7, 'max_iter': 200}
5 Adjusted Rand-Index: 0.152
6 Homogeneity Score: 0.504
7 Silhouette Score: 0.044
8 Normed Mutual-Info Score: 0.432

```



```

1 | 6.534716367721558 seconds.

```

- The above solution generated too many clusters to be properly visualized. However, the `Mutual_Info` score is quite decent because datapoints may be falling onto pockets that resemble the true labels.

5.3.5. Spectral Clustering CBOW

[Back to Outline](#)

- SpectralClustering can't be used with GridSearchCV because it lacks a `.fit` method. Therefore I won't use the function here.

```

1 clust= SpectralClustering()
2
3 params = {
4     'n_clusters':np.arange(10,26,5),
5     #'eigen_solver':['arpack','lobpcg',None],
6     'n_init':[15,25],
7     'assign_labels':['kmeans','discretize']
8 }
9
10 features='BOW'
11
12 i=3
13
14 t0=time()
15
16 y_pred = clust.fit_predict(X)
17
18 ari = adjusted_rand_score(y, y_pred)
19 performance.loc[i,'ARI'] = ari
20 print("Adjusted Rand-Index: %.3f" % ari)
21
22 hom = homogeneity_score(y,y_pred)
23 performance.loc[i,'Homogeneity'] = hom
24 print("Homogeneity Score: %.3f" % hom)
25
26 sil = silhouette_score(X,y_pred)
27 performance.loc[i,'Silhouette'] = sil
28 print("Silhouette Score: %.3f" % sil)
29
30 nmi = normalized_mutual_info_score(y,y_pred)
31 performance.loc[i,'Mutual_Info'] = nmi
32 print("Normed Mutual-Info Score: %.3f" % nmi)
33
34 performance.loc[i,'n_train'] = len(X)
35 performance.loc[i,'Features'] = features
36 performance.loc[i,'Algorithm'] = clust.__class__.__name__
37
38 # Print contingency matrix

```

```


39 crosstab = pd.crosstab(y, y_pred)
40 plt.figure(figsize=(8,5))
41 sns.heatmap(crosstab, annot=True,fmt='d', cmap=plt.cm.copper)
42 plt.show()
43 print(time()-t0,"seconds.")

```

```

1 Adjusted Rand-Index: 0.000
2 Homogeneity Score: 0.012
3 Silhouette Score: -0.344
4 Normed Mutual-Info Score: 0.063

```



```

1 0.6043310165405273 seconds.

```

```

1 performance.iloc[:, :7]

```

	Algorithm	n_train	Features	ARI	Homogeneity	Silhouette	Mutual_Info
0	KMeans	500	BOW	0.210114	0.446899	0.0498203	0.428272
1	MeanShift	500	BOW	0.00171957	0.101235	0.2766	0.17132
2	AffinityPropagation	500	BOW	0.151854	0.504396	0.0440665	0.431608
3	SpectralClustering	500	BOW	2.5954e-05	0.0120565	-0.344458	0.0632249

- Based on `Mutual_Info`, our highest score came from `AffinityPropagation`. However, the large number of clusters dividing our articles makes the solution a bit impractical.
- Fortunately we can perform supervised classification on this dataset because we actually do know who wrote these articles.

5.4. Classification on BOW

[Back to Outline](#)

5.4.1. Supervised Parameter Search Function

- The following function will print cross-validation, train and test accuracy scores in addition to the clustering scores we've been utilizing previously.
- The `GridSearchCV` will also find the parameters that produce the highest `Normalized Mutual Information` score.
- There is a very clear correlation between the `Mutual_Info` score and the `Test_Accuracy` from our classifiers.
- Notice that here the `n_train` will be smaller than in the previous section because here we are actually doing a train/test split, whereas in the previous section we used `fit_predict(X)` on the clustering algorithms.

```

1 def score_optimization(clf, params, features, i):
2     t0 = time()
3     # Heading
4     print('\n', '-'*40, '\n', clf.__class__.__name__, '\n', '-'*40)
5
6     # Find best parameters based on scoring of choice
7     score = make_scorer(normalized_mutual_info_score)
8     search = GridSearchCV(clf, params,
9                           scoring=score, cv=3).fit(X, y)
10    # Extract best estimator

```

```

11 best = search.best_estimator_
12 print("Best parameters:", search.best_params_)
13
14 # Cross-validate on all the data
15 cv = cross_val_score(X=X, y=y, estimator=best, cv=5)
16 print("\nCross-val scores(All Data):", cv)
17 print("Mean cv score:", cv.mean())
18 performance.loc[i, 'Cross_Val'] = cv.mean()
19
20 # Get train accuracy
21 best = best.fit(X_train, y_train)
22 train = best.score(X=X_train, y=y_train)
23 performance.loc[i, 'Train_Accuracy'] = train
24 print("\nTrain Accuracy Score:", train)
25
26 # Get test accuracy
27 test = best.score(X=X_test, y=y_test)
28 performance.loc[i, 'Test_Accuracy'] = test
29 print("\nTest Accuracy Score:", test)
30
31 y_pred = best.predict(X_test)
32
33 ari = adjusted_rand_score(y_test, y_pred)
34 performance.loc[i, 'ARI'] = ari
35 print("\nAdjusted Rand-Index: %.3f" % ari)
36
37 hom = homogeneity_score(y_test, y_pred)
38 performance.loc[i, 'Homogeneity'] = hom
39 print("Homogeneity Score: %.3f" % hom)
40
41 sil = silhouette_score(X_test, y_pred)
42 performance.loc[i, 'Silhouette'] = sil
43 print("Silhouette Score: %.3f" % sil)
44
45 nmi = normalized_mutual_info_score(y_test, y_pred)
46 performance.loc[i, 'Mutual_Info'] = nmi
47 print("Normed Mutual-Info Score: %.3f" % nmi)
48
49 #print(classification_report(y_test, y_pred))
50
51 conf_matrix = pd.crosstab(y_test, y_pred)
52 sns.heatmap(conf_matrix, annot=True, fmt='d', cmap=plt.cm.copper)
53 plt.show()
54
55 performance.loc[i, 'n_train'] = len(X_train)
56 performance.loc[i, 'Features'] = features
57 performance.loc[i, 'Algorithm'] = clf.__class__.__name__
58 print(time()-t0, 'seconds.')

```

5.4.2. Logistic Regression CBOW

[Back to Outline](#)

```

1 # Parameters to optimize
2 params = [{
3     'solver': ['newton-cg', 'lbfgs', 'sag'],
4     'C': [0.3, 0.5, 0.7, 1],
5     'penalty': ['l2']
6 }, {
7     'solver': ['liblinear', 'saga'],
8     'C': [0.3, 0.5, 0.7, 1],
9     'penalty': ['l1', 'l2']
10 }]
11
12 clf = LogisticRegression(
13     n_jobs=-1 # Use all CPU
14 )
15
16 score_optimization(clf=clf, params=params, features='BOW', i=4)

```

```


1 -----
2 LogisticRegression
3 -----
4 Best parameters: {'C': 0.7, 'penalty': 'l2', 'solver': 'liblinear'}
5
6 Cross-val scores(All Data): [0.79 0.78 0.74 0.82 0.65]
7 Mean cv score: 0.756
8
9 Train Accuracy Score: 1.0

```

```

10
11 Test Accuracy Score: 0.8416666666666667
12
13 Adjusted Rand-Index: 0.720
14 Homogeneity Score: 0.834
15 Silhouette Score: -0.002
16 Normed Mutual-Info Score: 0.838

```



```

1 465.8297345638275 seconds.

```

- Although the clustering results didn't have a train/test or cross-validation score, here we have a `Mutual_Info` score around twice the highest of our clusters. Above, `Mutual_Info` was very close to `Accuracy`, just two percentage points away. As we get more solutions we'll see the consistency between `Mutual_Info` and `Accuracy` among other classifiers. This will allow us to assess classification and clustering solutions by a fair mutual metric.

5.4.3. Random Forest CBOW

[Back to Outline](#)

```

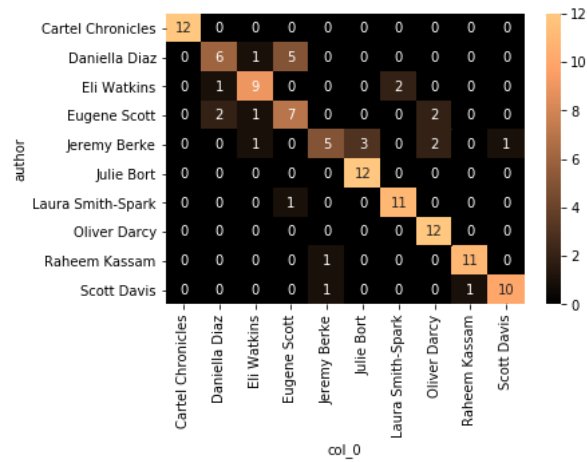
1 # Parameters to compare
2 params = {
3     'criterion':['entropy','gini'],
4 }
5
6 # Implement the classifier
7 clf = ensemble.RandomForestClassifier(
8     n_estimators=100,
9     max_features=None,
10    n_jobs=-1,
11 )
12
13 score_optimization(clf=clf,params=params,features='BOW',i=5)

```

```

1 -----
2 RandomForestClassifier
3 -----
4 Best parameters: {'criterion': 'gini'}
5
6 Cross-val scores(All Data): [0.78 0.82 0.73 0.83 0.7 ]
7 Mean cv score: 0.772
8
9 Train Accuracy Score: 1.0
10
11 Test Accuracy Score: 0.7916666666666666
12
13 Adjusted Rand-Index: 0.635
14 Homogeneity Score: 0.759
15 Silhouette Score: -0.050
16 Normed Mutual-Info Score: 0.763

```



1 | 32.785361528396606 seconds.

5.4.4. Gradient Boosting CBOW

[Back to Outline](#)

```

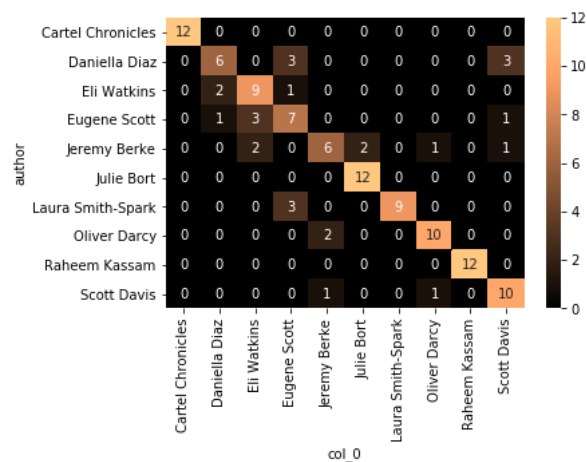
1 # Parameters to compare
2 params = {
3     'learning_rate':[0.3,0.5,0.7,1]
4 }
5
6 # Implement the classifier
7 clf = ensemble.GradientBoostingClassifier(
8     max_features=None
9 )
10
11 score_optimization(clf=clf,params=params,features='BOW',i=6)

```

```

1 -----
2 GradientBoostingClassifier
3 -----
4 Best parameters: {'learning_rate': 0.3}
5
6 Cross-val scores(All Data): [0.74 0.77 0.72 0.79 0.67]
7 Mean cv score: 0.738
8
9 Train Accuracy Score: 1.0
10
11 Test Accuracy Score: 0.775
12
13 Adjusted Rand-Index: 0.598
14 Homogeneity Score: 0.742
15 Silhouette Score: -0.022
16 Normed Mutual-Info Score: 0.745

```




```
1 158.82711482048035 seconds.
```

Results

- Clearly classifiers obtain higher scores than clustering, this is despite being trained with less data.
- So far `Accuracy` correlates perfectly with `Mutual_Info`.

```
1 performance.iloc[:7].sort_values('Mutual_Info', ascending=False)
  [['Algorithm', 'n_train', 'Features', 'Mutual_Info', 'Test_Accuracy']]
```

	Algorithm	n_train	Features	Mutual_Info	Test_Accuracy
4	LogisticRegression	380	BOW	0.837992	0.841667
5	RandomForestClassifier	380	BOW	0.762783	0.791667
6	GradientBoostingClassifier	380	BOW	0.744608	0.775
2	AffinityPropagation	500	BOW	0.431608	NaN
0	KMeans	500	BOW	0.428272	NaN
1	MeanShift	500	BOW	0.17132	NaN
3	SpectralClustering	500	BOW	0.0632249	NaN

6. Unsupervised Feature Generation

[Back to Outline](#)

6.1. Latent Semantic Analysis

- Different from Bag-of-Words, Latent Semantic Analysis doesn't identify the most common words present in each article. Instead it identifies thematic components present in the text. Each cell doesn't contain a count, but rather a measure of how well a given feature is exemplified by the current document.

```
1 vectorizer = TfidfVectorizer(max_df=0.3, # drop words that occur in more than X percent of documents
2                             min_df=8, # only use words that appear at least X times
3                             stop_words='english',
4                             lowercase=True, #convert everything to lower case
5                             use_idf=True, #we definitely want to use inverse document frequencies in our weighting
6                             norm='l2', #Applies a correction factor so that longer paragraphs and shorter
  paragraphs get treated equally
7                             smooth_idf=True #Adds 1 to all document frequencies, as if an extra document existed
  that used every word once. Prevents divide-by-zero errors
8                             )
9
10 #Pass pandas series to our vectorizer model
11 counts_tfidf = vectorizer.fit_transform(bow_counts.content)
12
13
```

- Notice that the content fed into the vectorizer is the same amount of data we used for BOW Counts. (500 articles in total, 50 by each author). We could use all of the 1000 articles, but first let's compare the LSA performance against BOW using the same data.
- The vectorizer returns a CSR Matrix which can then be reduced as in PCA.

```
1 counts_tfidf
```

```
1 <500x2537 sparse matrix of type '<class 'numpy.float64'>'
2   with 60411 stored elements in Compressed Sparse Row format>
```

- Reducing to 460 features will retain 98% of the explained variance.

```
1 svd = TruncatedSVD(460)
2 svd.fit(counts_tfidf)
3 svd.explained_variance_ratio_.sum()
```

```
1 | 0.9859640295175764
```

```
1 | lsa = make_pipeline(svd, Normalizer(copy=False))
2 | lsa_data = lsa.fit_transform(counts_tfidf)
3 | lsa_data.shape
```

```
1 | (500, 460)
```

```
1 | lsa_data = pd.DataFrame(lsa_data)
2 | lsa_data.head()
```

	0	1	2	3	4	5	6	7	8	9	...	450
0	0.134680	-0.089814	-0.126869	-0.076195	-0.029081	-0.048506	-0.178471	-0.246023	0.214692	-0.006315	...	-0.004...
1	0.114785	-0.061037	-0.066941	-0.013864	0.027526	-0.048437	-0.012254	-0.015127	-0.030769	0.007581	...	-0.000...
2	0.111855	-0.089993	-0.123386	-0.087870	-0.030889	-0.047380	-0.150957	-0.226742	0.189297	0.005570	...	-0.002...
3	0.095245	-0.089762	-0.090381	-0.076953	-0.007260	-0.075972	-0.205302	-0.271548	0.197548	0.007419	...	-0.019...
4	0.071403	-0.056780	-0.076383	-0.049186	-0.022184	-0.048335	-0.125046	-0.123988	0.148768	0.009617	...	-0.013...

5 rows × 460 columns

6.2. Clustering on LSA (BOW Content)

- We'll repeat the clustering and classification, now using the LSA features from the same 500 articles we used in BOW Counts.

[Back to Outline](#)

```
1 | #First, establish X and Y
2 | y = bow_counts['author']
3 | X = lsa_data
4 |
5 | X_train, X_test, y_train, y_test = train_test_split(X,
6 |                                                    y,
7 |                                                    test_size=0.24,
8 |                                                    random_state=0,
9 |                                                    stratify=y)
```

```
1 | y_test.value_counts()
```

```
1 | Raheem Kassam      12
2 | Julie Bort         12
3 | Eli Watkins        12
4 | Eugene Scott       12
5 | Cartel Chronicles  12
6 | Scott Davis        12
7 | Jeremy Berke       12
8 | Laura Smith-Spark  12
9 | Oliver Darcy       12
10 | Daniella Diaz      12
11 | Name: author, dtype: int64
```

6.2.2. KMeans LSA


[Back to Outline](#)

```
1 | clust=KMeans()
2 | params={
3 |     'n_clusters': np.arange(10,30,5),
4 |     'init': ['k-means++', 'random'],
5 |     'n_init': [10,20],
6 |     'precompute_distances': [True, False]
7 | }
8 | evaluate_clust(clust, params, features='LSA', i=7)
```

```

1 -----
2 KMeans
3 -----
4 Best parameters: {'init': 'random', 'n_clusters': 20, 'n_init': 10, 'precompute_distances': True}
5 Adjusted Rand-Index: 0.336
6 Homogeneity Score: 0.534
7 Silhouette Score: 0.048
8 Normed Mutual-Info Score: 0.482

```



```

1 27.749962329864502 seconds.

```

6.2.3. Mean Shift LSA

[Back to Outline](#)

```


1 #Declare and fit the model
2 clust = MeanShift()
3
4 params={
5     'bandwidth':[0.5,0.7,0.9]
6 }
7 evaluate_clust(clust,params,features='LSA',i=8)

```

```

1 -----
2 MeanShift
3 -----
4 Best parameters: {'bandwidth': 0.5}
5 Adjusted Rand-Index: 0.001
6 Homogeneity Score: 1.000
7 Silhouette Score: 0.010
8 Normed Mutual-Info Score: 0.609

```



```

1 16.05977153778076 seconds.


```

6.2.4. Affinity Propagation LSA

[Back to Outline](#)

```
1 #Declare and fit the model.
2 clust = AffinityPropagation()
3
4 params = {
5     'damping': [.5, .7, .9],
6     'max_iter': [200, 500]
7 }
8 evaluate_clust(clust, params, features='LSA', i=9)
```

```
1 -----
2 AffinityPropagation
3 -----
4 Best parameters: {'damping': 0.5, 'max_iter': 200}
5 Adjusted Rand-Index: 0.110
6 Homogeneity Score: 0.689
7 Silhouette Score: 0.063
8 Normed Mutual-Info Score: 0.507
```



```
1 3.9195096492767334 seconds.
```

6.2.5. Spectral Clustering LSA

[Back to Outline](#)

SpectralClustering can't be used with GridSearchCV because it lacks a .fit method. Therefore I won't use the function here.

```
1 clust= SpectralClustering()
2
3 params = {
4     'n_clusters': np.arange(10, 26, 5),
5     #'eigen_solver': ['arpack', 'lobpcg', None],
6     'n_init': [15, 25],
7     'assign_labels': ['kmeans', 'discretize']
8 }
9
10 features='LSA'
11
12 i=10
13
14 t0=time()
15
16 y_pred = clust.fit_predict(X)
17
18 ari = adjusted_rand_score(y, y_pred)
19 performance.loc[i, 'ARI'] = ari
20 print("Adjusted Rand-Index: %.3f" % ari)
21
22 hom = homogeneity_score(y, y_pred)
23 performance.loc[i, 'Homogeneity'] = hom
24 print("Homogeneity Score: %.3f" % hom)
25
26 sil = silhouette_score(X, y_pred)
27 performance.loc[i, 'Silhouette'] = sil
28 print("Silhouette Score: %.3f" % sil)
29
30 nmi = normalized_mutual_info_score(y, y_pred)
```

```


31 performance.loc[i,'Mutual_Info'] = nmi
32 print("Normed Mutual-Info Score: %.3f" % nmi)
33
34 performance.loc[i,'n_train'] = len(X)
35 performance.loc[i,'Features'] = features
36 performance.loc[i,'Algorithm'] = clust.__class__.__name__
37
38 # Print contingency matrix
39 crosstab = pd.crosstab(y, y_pred)
40 plt.figure(figsize=(8,5))
41 sns.heatmap(crosstab, annot=True,fmt='d', cmap=plt.cm.copper)
42 plt.show()
43 print(time()-t0,"seconds.")

```

```

1 Adjusted Rand-Index: 0.075
2 Homogeneity Score: 0.270
3 Silhouette Score: 0.040
4 Normed Mutual-Info Score: 0.369

```



```

1 0.46030735969543457 seconds.

```

Results (See below)

- Based on `Mutual_Info` score, classification outperforms clustering regardless of the method used for feature-generation.
- Within the clustering solutions however, LSA produced higher scores than BOW except for SpectralClustering.

```

1 performance.iloc[:11].sort_values('Mutual_Info',ascending=False)
  [['Algorithm', 'n_train', 'Features', 'Mutual_Info', 'Test_Accuracy']]

```

	Algorithm	n_train	Features	Mutual_Info	Test_Accuracy
4	LogisticRegression	380	BOW	0.837992	0.841667
5	RandomForestClassifier	380	BOW	0.762783	0.791667
6	GradientBoostingClassifier	380	BOW	0.744608	0.775
8	MeanShift	500	LSA	0.609241	NaN
9	AffinityPropagation	500	LSA	0.506508	NaN
7	KMeans	500	LSA	0.482019	NaN
2	AffinityPropagation	500	BOW	0.431608	NaN
0	KMeans	500	BOW	0.428272	NaN
10	SpectralClustering	500	LSA	0.368991	NaN
1	MeanShift	500	BOW	0.17132	NaN
3	SpectralClustering	500	BOW	0.0632249	NaN

6.3. Classification on LSA (BOW Content)

[Back to Outline](#)

- Now we'll do supervised classification on the LSA features.

6.3.1. Logistic Regression LSA

```

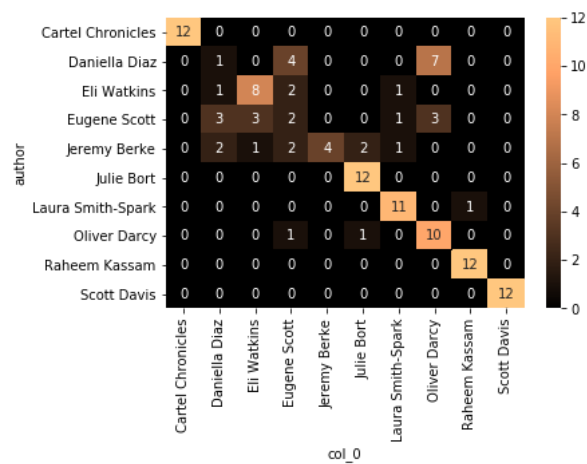
1 # Parameters to optimize
2 params = [{
3     'solver': ['newton-cg', 'lbfgs', 'sag'],
4     'C': [0.3, 0.5, 0.7, 1],
5     'penalty': ['l2']
6 }, {
7     'solver': ['liblinear', 'saga'],
8     'C': [0.3, 0.5, 0.7, 1],
9     'penalty': ['l1', 'l2']
10 }]
11
12 clf = LogisticRegression(
13     n_jobs=-1 # Use all CPU
14 )
15
16 score_optimization(clf=clf, params=params, features='LSA', i=11)

```

```

1 -----
2 LogisticRegression
3 -----
4 Best parameters: {'C': 1, 'penalty': 'l2', 'solver': 'liblinear'}
5
6 Cross-val scores(All Data): [0.7 0.63 0.71 0.72 0.6 ]
7 Mean cv score: 0.6719999999999999
8
9 Train Accuracy Score: 0.9789473684210527
10
11 Test Accuracy Score: 0.7
12
13 Adjusted Rand-Index: 0.585
14 Homogeneity Score: 0.714
15 Silhouette Score: 0.048
16 Normed Mutual-Info Score: 0.724

```



```

1 | 14.797184228897095 seconds.

```

6.3.2. Random Forest LSA

[Back to Outline](#)

```

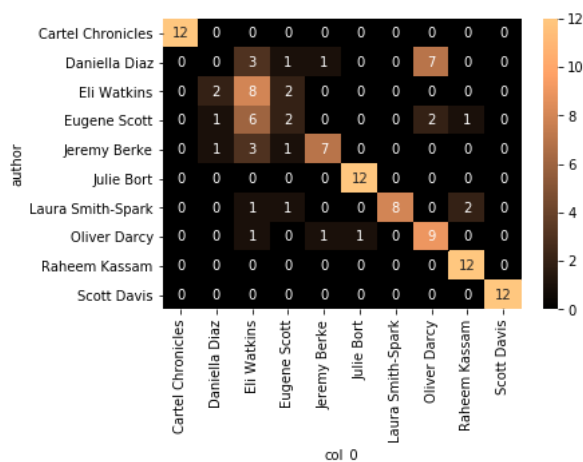
1 # Parameters to compare
2 params = {
3     'criterion':['entropy','gini'],
4 }
5
6 # Implement the classifier
7 clf = ensemble.RandomForestClassifier(
8     n_estimators=100,
9     max_features=None,
10    n_jobs=-1,
11 )
12
13 score_optimization(clf=clf,params=params,features='LSA',i=12)

```

```

1 -----
2 RandomForestClassifier
3 -----
4 Best parameters: {'criterion': 'entropy'}
5
6 Cross-val scores(All Data): [0.55 0.6  0.63 0.65 0.6 ]
7 Mean cv score: 0.606
8
9 Train Accuracy Score: 1.0
10
11 Test Accuracy Score: 0.6833333333333333
12
13 Adjusted Rand-Index: 0.538
14 Homogeneity Score: 0.691
15 Silhouette Score: 0.043
16 Normed Mutual-Info Score: 0.705

```



```

1 | 129.73267102241516 seconds.

```

6.3.3. Gradient Boosting LSA

[Back to Outline](#)

```

1 # Parameters to compare
2 params = {
3     'learning_rate':[0.3,0.5,0.7,1]
4 }
5
6 # Implement the classifier
7 clf = ensemble.GradientBoostingClassifier(
8     max_features=None
9 )
10
11 score_optimization(clf=clf,params=params,features='LSA',i=13)

```

```

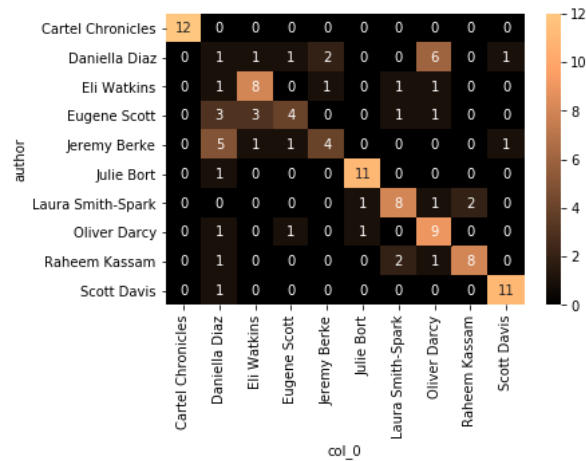
1 -----
2 GradientBoostingClassifier
3 -----
4 Best parameters: {'learning_rate': 0.3}
5
6 Cross-val scores(All Data): [0.56 0.5  0.54 0.58 0.52]

```

```

7 | Mean cv score: 0.54
8 |
9 | Train Accuracy Score: 1.0
10 |
11 | Test Accuracy Score: 0.6333333333333333
12 |
13 | Adjusted Rand-Index: 0.444
14 | Homogeneity Score: 0.602
15 | Silhouette Score: 0.042
16 | Normed Mutual-Info Score: 0.607

```



```

1 | 108.13380837440491 seconds.

```

Results

- Once again, classification trumps clustering regardless of the feature-generation method.
- BOW features have performed consistently better than LSA on all classifiers.

```

1 | performance.iloc[:14].sort_values('Mutual_Info',ascending=False)
  | [['Algorithm', 'n_train', 'Features', 'Mutual_Info', 'Test_Accuracy']].iloc[:9]

```

	Algorithm	n_train	Features	Mutual_Info	Test_Accuracy
4	LogisticRegression	380	BOW	0.837992	0.841667
5	RandomForestClassifier	380	BOW	0.762783	0.791667
6	GradientBoostingClassifier	380	BOW	0.744608	0.775
11	LogisticRegression	380	LSA	0.724234	0.7
12	RandomForestClassifier	380	LSA	0.705238	0.683333
8	MeanShift	500	LSA	0.609241	NaN
13	GradientBoostingClassifier	380	LSA	0.607244	0.633333
9	AffinityPropagation	500	LSA	0.506508	NaN
7	KMeans	500	LSA	0.482019	NaN

6.4. Clustering on LSA (All Content)

[Back to Outline](#)

- Since LSA allows for very quick feature-generation, it's worth making a comparison between past results VS the utilization of all available data. After all, the LSA classifiers aren't far behind the BOW classifiers on 380 samples. With twice the number of articles LSA could very well outperform BOW.


```

1 vectorizer = TfidfVectorizer(max_df=0.3, # drop words that occur in more than X percent of documents
2                             min_df=8, # only use words that appear at least X times
3                             stop_words='english',
4                             lowercase=True, #convert everything to lower case
5                             use_idf=True, #we definitely want to use inverse document frequencies in our weighting
6                             norm='l2', #Applies a correction factor so that longer paragraphs and shorter
    paragraphs get treated equally
7                             smooth_idf=True #Adds 1 to all document frequencies, as if an extra document existed
    that used every word once. Prevents divide-by-zero errors
8                             )
9
10 #Pass pandas series to our vectorizer model
11 counts_tfidf = vectorizer.fit_transform(authors_data.content)
12
13

```

- Notice that this time we fed all the articles into the vectorizer. See the size of the CSR Matrix underneath. The 1000 rows are 100 articles for each 10 authors.

```

1 counts_tfidf

```

```

1 <1000x4141 sparse matrix of type '<class 'numpy.float64''>'
2 with 129747 stored elements in Compressed Sparse Row format>

```

- This time we need 900 features to retain 98% of the variance.

```

1 svd = TruncatedSVD(900)
2 svd.fit(counts_tfidf)
3 svd.explained_variance_ratio_.sum()

```

```

1 0.983515568253816

```

```

1 lsa = make_pipeline(svd, Normalizer(copy=False))
2 lsa_data = lsa.fit_transform(counts_tfidf)
3 lsa_data.shape

```

```

1 (1000, 900)

```

```

1 lsa_data = pd.DataFrame(lsa_data)
2 lsa_data.head()

```

	0	1	2	3	4	5	6	7	8	9	...	890
0	0.095554	-0.077466	-0.088685	0.116715	-0.029071	0.229682	-0.093686	0.058752	0.031923	0.028790	...	-0.0105
1	0.077584	-0.042902	-0.065056	0.044378	0.013497	-0.008408	-0.020459	-0.026143	-0.014150	-0.002720	...	-0.0117
2	0.097580	-0.087290	-0.111300	0.132449	-0.040793	0.276451	-0.112235	0.081090	0.037374	0.025789	...	-0.0055
3	0.078344	-0.081317	-0.090485	0.133814	-0.014822	0.324981	-0.154498	0.116849	0.007485	0.040119	...	-0.0023
4	0.054085	-0.044813	-0.057238	0.071518	-0.018558	0.141139	-0.072288	0.014504	0.024686	0.022923	...	0.0135

5 rows × 900 columns

```

1 #First, establish X and Y
2 y = authors_data['author']
3 X = lsa_data
4
5 X_train, X_test, y_train, y_test = train_test_split(X,
6                                                     y,
7                                                     test_size=0.24,
8                                                     random_state=0,
9                                                     stratify=y)

```

- The test data reflects the change in size.

```

1 y_test.value_counts()

```

```

1 Raheem Kassam      24
2 Cartel Chronicles  24
3 Scott Davis        24
4 Jeremy Berke       24
5 Daniella Diaz      24
6 Eli Watkins        24
7 Eugene Scott       24
8 Julie Bort         24
9 Oliver Darcy       24
10 Laura Smith-Spark  24
11 Name: author, dtype: int64

```

6.4.1. KMeans LSA (All Content)

[Back to Outline](#)

```


1 clust=KMeans()
2 params={
3     'n_clusters': np.arange(10,30,5),
4     'init': ['k-means++','random'],
5     'n_init':[10,20],
6     'precompute_distances':[True,False]
7 }
8 evaluate_clust(clust,params,features='LSA',i=14)

```

```

1 -----
2 KMeans
3 -----
4 Best parameters: {'init': 'random', 'n_clusters': 10, 'n_init': 10, 'precompute_distances': True}
5 Adjusted Rand-Index: 0.246
6 Homogeneity Score: 0.463
7 Silhouette Score: 0.017
8 Normed Mutual-Info Score: 0.502

```



```

1 133.5461344718933 seconds.

```

6.4.2. Mean Shift LSA (All Content)

[Back to Outline](#)

```


1 #Declare and fit the model
2 clust = MeanShift()
3
4 params={
5     'bandwidth':[0.5,0.7,0.9]
6 }
7 evaluate_clust(clust,params,features='LSA',i=15)

```

```

1 -----
2 MeanShift
3 -----
4 Best parameters: {'bandwidth': 0.5}
5 Adjusted Rand-Index: 0.000
6 Homogeneity Score: 0.999
7 Silhouette Score: 0.004
8 Normed Mutual-Info Score: 0.577

```



```

1 60.95715284347534 seconds.

```

6.4.3. Affinity Propagation LSA (All Content)

[Back to Outline](#)

```


1 #Declare and fit the model.
2 clust = AffinityPropagation()
3
4 params = {
5     'damping': [.5, .7, .9],
6     'max_iter': [200, 500]
7 }
8 evaluate_clust(clust, params, features='LSA', i=16)

```

```

1 -----
2 AffinityPropagation
3 -----
4 Best parameters: {'damping': 0.5, 'max_iter': 200}
5 Adjusted Rand-Index: 0.072
6 Homogeneity Score: 0.679
7 Silhouette Score: 0.055
8 Normed Mutual-Info Score: 0.467

```



```

1 18.127511978149414 seconds.

```


6.4.4. Spectral Clustering LSA (All Content)

[Back to Outline](#)

SpectralClustering can't be used with GridSearchCV because it lacks a .fit method. Therefore I won't use the function here.

```
1 clust= SpectralClustering()
2
3 params = {
4     'n_clusters':np.arange(10,26,5),
5     #'eigen_solver':['arpack','lobpcg',None],
6     'n_init':[15,25],
7     'assign_labels':['kmeans','discretize']
8 }
9
10 features='LSA'
11
12 i=17
13
14 t0=time()
15
16 y_pred = clust.fit_predict(X)
17
18 ari = adjusted_rand_score(y, y_pred)
19 performance.loc[i,'ARI'] = ari
20 print("Adjusted Rand-Index: %.3f" % ari)
21
22 hom = homogeneity_score(y,y_pred)
23 performance.loc[i,'Homogeneity'] = hom
24 print("Homogeneity Score: %.3f" % hom)
25
26 sil = silhouette_score(X,y_pred)
27 performance.loc[i,'Silhouette'] = sil
28 print("Silhouette Score: %.3f" % sil)
29
30 nmi = normalized_mutual_info_score(y,y_pred)
31 performance.loc[i,'Mutual_Info'] = nmi
32 print("Normed Mutual-Info Score: %.3f" % nmi)
33
34 performance.loc[i,'n_train'] = len(X)
35 performance.loc[i,'Features'] = features
36 performance.loc[i,'Algorithm'] = clust.__class__.__name__
37
38 # Print contingency matrix
39 crosstab = pd.crosstab(y, y_pred)
40 plt.figure(figsize=(8,5))
41 sns.heatmap(crosstab, annot=True,fmt='d', cmap=plt.cm.copper)
42 plt.show()
43 print(time()-t0,"seconds.")
```

```
1 Adjusted Rand-Index: 0.125
2 Homogeneity Score: 0.335
3 Silhouette Score: 0.031
4 Normed Mutual-Info Score: 0.417
```



```
1 0.7042851448059082 seconds.
```

6.5. Classification on LSA (All Content)

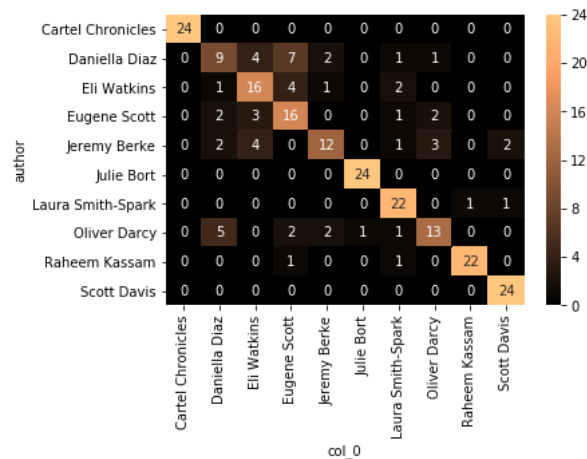
[Back to Outline](#)

- We've done clustering on LSA using all 1000 articles. Now let's classify.

6.5.1. Logistic Regression LSA (All Content)

```
1 # Parameters to optimize
2 params = [{
3     'solver': ['newton-cg', 'lbfgs', 'sag'],
4     'C': [0.3, 0.5, 0.7, 1],
5     'penalty': ['l2']
6 }, {
7     'solver': ['liblinear', 'saga'],
8     'C': [0.3, 0.5, 0.7, 1],
9     'penalty': ['l1', 'l2']
10 }]
11
12 clf = LogisticRegression(
13     n_jobs=-1 # Use all CPU
14 )
15
16 score_optimization(clf=clf, params=params, features='LSA', i=18)
```

```
1 -----
2 LogisticRegression
3 -----
4 Best parameters: {'C': 1, 'penalty': 'l2', 'solver': 'liblinear'}
5
6 Cross-val scores(All Data): [0.66 0.68 0.705 0.695 0.62 ]
7 Mean cv score: 0.6719999999999999
8
9 Train Accuracy Score: 0.9631578947368421
10
11 Test Accuracy Score: 0.7583333333333333
12
13 Adjusted Rand-Index: 0.592
14 Homogeneity Score: 0.685
15 Silhouette Score: 0.030
16 Normed Mutual-Info Score: 0.687
```



```
1 | 56.86849617958069 seconds.
```

6.5.2. Random Forest LSA

[Back to Outline](#)

```

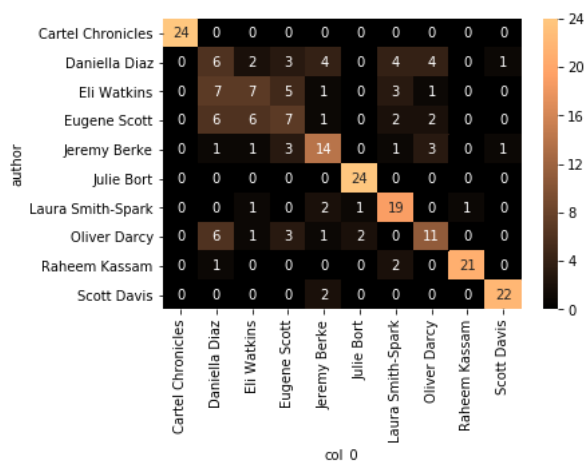
1 # Parameters to compare
2 params = {
3     'criterion':['entropy','gini'],
4 }
5
6 # Implement the classifier
7 clf = ensemble.RandomForestClassifier(
8     n_estimators=100,
9     max_features=None,
10    n_jobs=-1,
11 )
12
13 score_optimization(clf=clf,params=params,features='LSA',i=19)

```

```

1 -----
2 RandomForestClassifier
3 -----
4 Best parameters: {'criterion': 'entropy'}
5
6 Cross-val scores(All Data): [0.58  0.615 0.615 0.65  0.555]
7 Mean cv score: 0.603
8
9 Train Accuracy Score: 1.0
10
11 Test Accuracy Score: 0.6458333333333334
12
13 Adjusted Rand-Index: 0.481
14 Homogeneity Score: 0.589
15 Silhouette Score: 0.027
16 Normed Mutual-Info Score: 0.591

```



```

1 | 573.7529637813568 seconds.

```

6.5.3. Gradient Boosting LSA

[Back to Outline](#)

```

1 # Parameters to compare
2 params = {
3     'learning_rate':[0.3,0.5,0.7,1]
4 }
5
6 # Implement the classifier
7 clf = ensemble.GradientBoostingClassifier(
8     max_features=None
9 )
10
11 score_optimization(clf=clf,params=params,features='LSA',i=20)

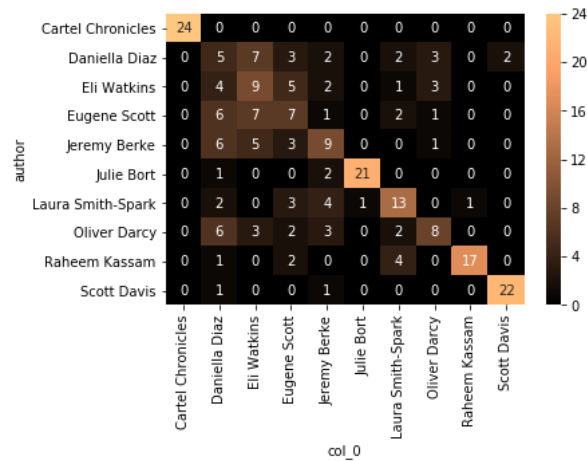
```

```

1 -----
2 GradientBoostingClassifier
3 -----
4 Best parameters: {'learning_rate': 0.3}
5
6 Cross-val scores(All Data): [0.56  0.545 0.54  0.585 0.525]

```

```
7 | Mean cv score: 0.5509999999999999
8 |
9 | Train Accuracy Score: 1.0
10 |
11 | Test Accuracy Score: 0.5625
12 |
13 | Adjusted Rand-Index: 0.371
14 | Homogeneity Score: 0.510
15 | Silhouette Score: 0.026
16 | Normed Mutual-Info Score: 0.512
```



```
1 | 581.0574629306793 seconds.
```

Comparing Results:

The results of more data are mixed with other methods. The LogisticRegression LSA with 760 samples is above GradientBoosting with 380, but below LogisticRegression with 380.

- **n_train.** Overall the 380 train size which is the 75% train split from the 500 BOW set generated higher scores than larger sizes.
- **Features.** Overall BOW features produced higher scores than most LSA features.
- **Supervised VS Unsupervised.** Classification produced indisputably higher scores than clustering regardless of size or feature-generation .

```
1 | performance.sort_values('Mutual_Info',ascending=False)
   | [['Algorithm','n_train','Features','Mutual_Info','Test_Accuracy']].iloc[:10]
```

	Algorithm	n_train	Features	Mutual_Info	Test_Accuracy
4	LogisticRegression	380	BOW	0.837992	0.841667
5	RandomForestClassifier	380	BOW	0.762783	0.791667
6	GradientBoostingClassifier	380	BOW	0.744608	0.775
11	LogisticRegression	380	LSA	0.724234	0.7
12	RandomForestClassifier	380	LSA	0.705238	0.683333
18	LogisticRegression	760	LSA	0.687342	0.758333
8	MeanShift	500	LSA	0.609241	NaN
13	GradientBoostingClassifier	380	LSA	0.607244	0.633333
19	RandomForestClassifier	760	LSA	0.590879	0.645833
15	MeanShift	1000	LSA	0.577176	NaN

7. Choosing Model

[Back to Outline](#)

7.1. Comparing Scores

- Since we tracked several scores throughout our testing, let's first compare our scores. The table below is sorted by `Mutual_Info` score.
- The first three scores `Mutual_Info`, `ARI`, `Homogeneity` are most commonly used for clustering. `Cross_Val`, `Train_Accuracy`, `Test_Accuracy` are limited to classification. Therefore the `NaN` missing values are the clustering algorithms.
- Notice that `Mutal_Info` scores and `Test_Accuracy` are very closely related to one another.
 - `Homogeneity` is close as well, but it gives 0.99 for index 15, which is a clustering algorithm with several dozens of clusters. `Homogeneity` will reward clustering solutions with numerous `n_clusters` because it penalizes clusters containing mixed `true_labels`. But so many clusters are practically useless.

```
1 performance.sort_values('Mutual_Info',ascending=False)
   [['Mutual_Info','ARI','Homogeneity','Cross_Val','Train_Accuracy','Test_Accuracy']].iloc[:10]
```

	Mutual_Info	ARI	Homogeneity	Cross_Val	Train_Accuracy	Test_Accuracy
4	0.837992	0.720231	0.834389	0.756	1	0.841667
5	0.762783	0.634581	0.759031	0.772	1	0.791667
6	0.744608	0.597551	0.741836	0.738	1	0.775
11	0.724234	0.584667	0.714154	0.672	0.978947	0.7
12	0.705238	0.538326	0.69103	0.606	1	0.683333
18	0.687342	0.592213	0.684947	0.672	0.963158	0.758333
8	0.609241	0.000588778	1	NaN	NaN	NaN
13	0.607244	0.444004	0.602122	0.54	1	0.633333
19	0.590879	0.480719	0.589479	0.603	1	0.645833
15	0.577176	6.87927e-05	0.999398	NaN	NaN	NaN

7.2. Sorting by Test_Accuracy

[Back to Outline](#)

Although we established that `Mutual_Info` and `Test_Accuracy` are very closely aligned, if we sort by `Test_Accuracy` there is a slight difference in top performers.

- First of all, since clustering solutions have missing values they are all at the bottom.
- All the BOW solutions are still at the top.
- LogisticRegression 760 LSA is now above itself at 380 samples. For RandomForest however, less samples produced a higher test accuracy. Same goes for GradientBoosting underneath.

```
1 performance.sort_values('Test_Accuracy',ascending=False)
   [['Algorithm','n_train','Features','Mutual_Info','Test_Accuracy']].iloc[:10]
```

	Algorithm	n_train	Features	Mutual_Info	Test_Accuracy
4	LogisticRegression	380	BOW	0.837992	0.841667
5	RandomForestClassifier	380	BOW	0.762783	0.791667
6	GradientBoostingClassifier	380	BOW	0.744608	0.775
18	LogisticRegression	760	LSA	0.687342	0.758333
11	LogisticRegression	380	LSA	0.724234	0.7
12	RandomForestClassifier	380	LSA	0.705238	0.683333
19	RandomForestClassifier	760	LSA	0.590879	0.645833
13	GradientBoostingClassifier	380	LSA	0.607244	0.633333
20	GradientBoostingClassifier	760	LSA	0.512496	0.5625
0	KMeans	500	BOW	0.428272	NaN

7.3. Winner

[Back to Outline](#)

Algorithm

Clearly LogisticRegression has done a better job at predicting the author name regardless of other factors. Across varying train size and feature-generation, 3 out of the 5 top solutions are from LogisticRegression.

Feature-Generation

For the purposes of predicting author's names, classification on BOW features has outperformed LSA. However LSA could be more appropriate for other tasks. Perhaps an author's uniqueness is more palpable from his vocabulary than from the semantics of his writing. This may explain why BOW was superior in this project.

Train_Size

Train size produced dubious variations in LSA. More data helped LogisticRegression but made others less accurate. It would be nice to see the effects of Train size in BOW, but that takes a long time.

Score

Normalized Mutual Information is definitely the best score with which to compare clustering and classification algorithms. Other scores also have a close resemblance, therefore I'd recommend to always compare several clustering scores.

```
1 plot_data = performance.sort_values('Test_Accuracy',ascending=False)
2 [[ 'Algorithm', 'Features', 'n_train', 'Test_Accuracy' ]].iloc[:9]
3
4 plot_data.n_train = plot_data.n_train.apply(lambda x: str(x))
5 plot_data['method'] = plot_data.Features+'_'+plot_data.n_train
6 %matplotlib inline
7
8 sns.catplot(col='Algorithm',x='method',y='Test_Accuracy',
9             data=plot_data,kind='bar')
10 plt.show()
```

